

Kill the Quote Spreadsheet

Reading the quotes turned out to be the easy part. The product is a comparison a buyer can defend, and the hard problem inside it is deciding whether two vendors are even quoting the same thing.

5 formats Excel, letterhead PDF, Word, angled photo and a one-line email, read live	133/133 prices within 1% of the answer key, after reading, matching and conversion	106 → 14 flagged prices collapsed into the decisions a buyer actually has to make	₹12.9 L a year left on the table by one expired ISO certificate, surfaced on the award screen
--	---	--	--

THE BET

I spent the time where buyers lose it: reading, normalising and trusting quotes. Drafting is deliberately thin; any model can write an RFQ. One rule runs through every screen: **AI reads and interprets, code calculates, the buyer decides**. Corrugated packaging was the category because its units are genuinely messy: per box, per kg, per 100, per 1,000, inches against millimetres.

FIVE DECISIONS, AND WHAT EACH ONE COST

- 1 Extraction copies; it never converts.
The model records each price exactly as written, with the vendor's own words as evidence. Plain code turns it into ₹ per piece, delivered, before GST, and every step is one click from the number: ₹49.02/kg → remove 18% GST → less 3% footnote discount → × 0.566 kg box = ₹22.81. **Cost:** two stages instead of one clever prompt. **Gain:** a misreading and a miscalculation can never look alike.
- 2 Uncertainty is a state on screen, not a footnote.
Confidence comes from evidence (source found, handwriting, size gap, more than 40% off last year), never from the model rating itself. Compare stays locked until every flag is approved, corrected, excluded or sent back to the vendor in an email the AI drafts from their own words. "Not quoted" is never zero, so a 27-of-30 vendor can't win on total. **Cost:** 14 decisions before the buyer sees a winner.
- 3 A near size is not the same box.
Match by our code, then by size and ply, then let the model read descriptions like "the namkeen 150g shipper"; rules veto any ply or product-type mismatch. Up to 5% off matches quietly, 5–10% matches flagged red, beyond 10% is a different box. A new vendor's "same as before" is a risk, not the incumbent's price.
- 4 Quality is judged by evidence, and its price is shown.
A certificate is judged by its date, not by the vendor's "yes". Shree Balaji looks 17% dearer than last year until GST and a 6-point footnote are applied; then it would win 19 of 30 lines against the cleared vendors, but its ISO lapsed two weeks before the due date. The award says both: **₹4.08 Cr with cleared vendors; ₹3.95 Cr if Balaji renews**. Ask, don't reject.
- 5 A fixed pipeline with an answer key, not an autonomous agent.
Extraction has to be explained to an auditor, so it is a pipeline scored against ground truth. Open-ended reasoning lives in the analyst, which answers by calling the same functions that draw the grid, plus a calculator. The chat and the screen cannot disagree, and every answer ends with how it got there.

EVIDENCE, NOT ADJECTIVES

Prices within 1% of the true ₹/pc	133 / 133
Quoted vs not-quoted, every vendor × line	150 / 150
Must-catch facts (footnote, cert date, handwriting, USD)	21 / 21
Word-only matching, repeat runs	4 / 4

Edit ₹42 to ₹39 in a vendor's email and every price, total and answer moves with it. **The miss that mattered:** once in four reads, one vendor silently shrank from 30 lines to 9. I made failure loud, not just rarer: unmatched items now reach the buyer as a risk.

DELIBERATELY LEFT OUT

Real email, portal, logins. Plumbing the brief let me stub; the AI loops are live.

Payment terms in landed cost. Needs each buyer's cost of funds; the analyst says so.

Negotiation, L1-match, POs. They follow the decision this tool makes defensible.

Fine-tuning. Five documents aren't training data; corrections are logged for later.

Live-formula Excel, mobile approval. Next, for finance and the MD.

THE BETTER PROBLEM

Is Gupta's 15 × 11 × 6-inch box our CB-309, when it is 9% taller and names no paper grade?

Is a lighter 18 BF "value-engineered" box the same box for less, or one that crushes in the monsoon? No OCR answers that, and it is where savings leak back as damaged goods. Next I'd build a spec-equivalence layer, learned from every buyer's accept and reject, so each RFQ makes the next one sharper.